In the Model Context Protocol (MCP), a server makes requests to a client primarily for **Sampling**. This allows the server (or its tools) to offload complex tasks back to the AI model managed by the client, without the server needing direct access to API keys or reasoning capabilities. [1, 2]

The main reasons a server requests information or actions from a client include:

- **Sampling (AI Completions):** A server can request the client's LLM to generate text, translate data, or evaluate a result. For example, if a travel booking server retrieves a list of flights, it might send the list to the client's LLM with a prompt asking, *"Based on the user's preferences, which flight is the best?"* [1, 2, 3, 4, 5]

- **Information Elicitation:** If a server needs more context from the user to fulfill a request, it can ask the client to prompt the user for clarification, additional parameters, or consent. [1]

- **Tool Discovery and State Notifications:** While not a direct functional request, servers notify clients when their toolsets (available actions) change. The client can then respond by making a subsequent request to pull the newly updated tool list. [1]

By keeping the communication two-way, the client retains full control over user permissions, security, and model usage. [1, 2]

StreamableHTTP for this cases —> SSE